

Practical Machine Learning Project

Sitaram Gautam

Feb 23, 2016

Background

This analysis corresponds to the prediction assignment for Practical Machine Learning course offered by John Hopkin's University through Coursera. The dataset for the assignment is obtained from <http://groupware.le>. Based on the website, the data is gathered using different sensors while 6 participants exercised in different fashions. The way they exercised are categorized into A, B, C, D, and E; A being the correct way of doing the exercise and B-E including some level of mistakes. The goal of the assignment is to build an algorithm that takes different fields in the dataset and predicts the associated exercise category correctly.

```
suppressMessages(library(caret))
suppressMessages(library(rpart))
suppressMessages(library(xgboost))

set.seed(977)

training_original <- read.csv(
  "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv",
  na.strings = c("NA", "#DIV/0!"),
  sep = ",",
  header = TRUE
)
testing_original <- read.csv(
  "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv",
  na.strings = c("NA", "#DIV/0!"),
  sep = ",",
  header = TRUE
)
```

Data cleaning and preprocessing For this analysis, I am going to use 3 data cleaning/preprocessing techniques which are described below:

1. Take only the fields related to exercises which are belt, arm, dumbbell, and forearm

```
# Take only the fields related to exercises which are belt, arm, dumbbell, and forearm
sensorColumns <-
  grep(pattern = "_belt|_arm|_dumbbell|_forearm", names(training_original))
```

```
# take only the columns associated to
training_original <- training_original[, c(sensorColumns, 160)]
```

2. Remove columns with NA

```
# remove columns with NA
cols_wo_na <- colSums(is.na(training_original)) == 0
training_original <- training_original[, cols_wo_na]
```

3. Check for features's variance

```
## check to see if there is any fields with near zero variability.
near_zero_vars <- nearZeroVar(training_original, saveMetrics = TRUE)
near_zero_vars
```

##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.101904	6.7781062	FALSE	FALSE
## pitch_belt	1.036082	9.3772296	FALSE	FALSE
## yaw_belt	1.058480	9.9734991	FALSE	FALSE
## total_accel_belt	1.063160	0.1477933	FALSE	FALSE
## gyros_belt_x	1.058651	0.7134849	FALSE	FALSE
## gyros_belt_y	1.144000	0.3516461	FALSE	FALSE
## gyros_belt_z	1.066214	0.8612782	FALSE	FALSE
## accel_belt_x	1.055412	0.8357966	FALSE	FALSE
## accel_belt_y	1.113725	0.7287738	FALSE	FALSE
## accel_belt_z	1.078767	1.5237998	FALSE	FALSE
## magnet_belt_x	1.090141	1.6664968	FALSE	FALSE
## magnet_belt_y	1.099688	1.5187035	FALSE	FALSE
## magnet_belt_z	1.006369	2.3290184	FALSE	FALSE
## roll_arm	52.338462	13.5256345	FALSE	FALSE
## pitch_arm	87.256410	15.7323412	FALSE	FALSE
## yaw_arm	33.029126	14.6570176	FALSE	FALSE
## total_accel_arm	1.024526	0.3363572	FALSE	FALSE
## gyros_arm_x	1.015504	3.2769341	FALSE	FALSE
## gyros_arm_y	1.454369	1.9162165	FALSE	FALSE
## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_dumbbell	1.022388	84.2065029	FALSE	FALSE
## pitch_dumbbell	2.277372	81.7449801	FALSE	FALSE
## yaw_dumbbell	1.132231	83.4828254	FALSE	FALSE
## total_accel_dumbbell	1.072634	0.2191418	FALSE	FALSE
## gyros_dumbbell_x	1.003268	1.2282132	FALSE	FALSE
## gyros_dumbbell_y	1.264957	1.4167771	FALSE	FALSE
## gyros_dumbbell_z	1.060100	1.0498420	FALSE	FALSE
## accel_dumbbell_x	1.018018	2.1659362	FALSE	FALSE
## accel_dumbbell_y	1.053061	2.3748853	FALSE	FALSE
## accel_dumbbell_z	1.133333	2.0894914	FALSE	FALSE
## magnet_dumbbell_x	1.098266	5.7486495	FALSE	FALSE
## magnet_dumbbell_y	1.197740	4.3012945	FALSE	FALSE
## magnet_dumbbell_z	1.020833	3.4451126	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE

```
## accel_forearm_y      1.059406      5.1116094    FALSE FALSE
## accel_forearm_z      1.006250      2.9558659    FALSE FALSE
## magnet_forearm_x      1.012346      7.7667924    FALSE FALSE
## magnet_forearm_y      1.246914      9.5403119    FALSE FALSE
## magnet_forearm_z      1.000000      8.5771073    FALSE FALSE
## classe                1.469581      0.0254816    FALSE FALSE
```

Columns nzv and zeroVar indicate that none of the features left have near-zero/zero variance i.e no need of features to be filtered out.

Data dividing into training and test set 60 % of the original training set is taken as training set for building model and rest as testing set

```
## Splitting the original training dataset into another training and a testing set.
inTrain <-
  createDataPartition(y = training_original$classe, p = 0.60, list = FALSE)

training <- training_original[inTrain,]
testing <- training_original[-inTrain,]
```

```
model_fit_dt <- rpart(classe ~ ., data = training, method = "class")
model_fit_dt
```

Building model

```
## n= 11776
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##      1) root 11776 8428 A (0.28 0.19 0.17 0.16 0.18)
##      2) roll_belt< 130.5 10798 7458 A (0.31 0.21 0.19 0.18 0.11)
##      4) pitch_forearm< -33.95 906      7 A (0.99 0.0077 0 0 0) *
##      5) pitch_forearm>=-33.95 9892 7451 A (0.25 0.23 0.21 0.2 0.12)
##      10) magnet_dumbbell_y< 439.5 8385 5996 A (0.28 0.18 0.24 0.19 0.11)
##      20) roll_forearm< 122.5 5214 3082 A (0.41 0.18 0.19 0.17 0.058)
##      40) magnet_dumbbell_z< -25.5 1802  573 A (0.68 0.2 0.019 0.077 0.025)
##      80) roll_forearm>=-136.5 1510  309 A (0.8 0.16 0.019 0.025 0.0046) *
##      81) roll_forearm< -136.5 292  173 B (0.096 0.41 0.017 0.35 0.13) *
##      41) magnet_dumbbell_z>=-25.5 3412 2479 C (0.26 0.17 0.27 0.22 0.076)
##      82) yaw_belt>=169.5 389   47 A (0.88 0.051 0 0.067 0.0026) *
##      83) yaw_belt< 169.5 3023 2090 C (0.19 0.19 0.31 0.23 0.085)
##      166) accel_dumbbell_y>=-40.5 2622 1927 D (0.21 0.21 0.22 0.27 0.092)
##      332) pitch_belt< -42.95 313   56 B (0.026 0.82 0.099 0.026 0.029) *
##      333) pitch_belt>=-42.95 2309 1622 D (0.24 0.12 0.24 0.3 0.1)
##      666) roll_belt>=125.5 571  235 C (0.36 0.03 0.59 0.014 0.0035)
##      1332) magnet_belt_z< -326.5 173    3 A (0.98 0 0.0058 0 0.012) *
##      1333) magnet_belt_z>=-326.5 398   63 C (0.095 0.043 0.84 0.02 0) *
##      667) roll_belt< 125.5 1738 1059 D (0.2 0.15 0.13 0.39 0.13)
```

```

##          1334) yaw_belt< -87.55 882  640 B (0.25 0.27 0.14 0.19 0.14)
##          2668) yaw_belt>=-93.25 732  490 B (0.29 0.33 0.16 0.061 0.15)
##          5336) magnet_forearm_z>=-43.5 330  135 A (0.59 0.2 0.058 0.067 0.088) *
##          5337) magnet_forearm_z< -43.5 402  225 B (0.045 0.44 0.25 0.057 0.21) *
##          2669) yaw_belt< -93.25 150   28 D (0.067 0 0.04 0.81 0.08) *
##          1335) yaw_belt>=-87.55 856  344 D (0.14 0.03 0.11 0.6 0.12)
##          2670) yaw_arm< -102.85 92    0 A (1 0 0 0 0) *
##          2671) yaw_arm>=-102.85 764  252 D (0.034 0.034 0.13 0.67 0.14) *
##          167) accel_dumbbell_y< -40.5 401  56 C (0.01 0.05 0.86 0.035 0.045) *
## 21) roll_forearm>=122.5 3171 2138 C (0.081 0.18 0.33 0.23 0.18)
##          42) accel_forearm_x>=-108.5 2255 1425 C (0.088 0.22 0.37 0.11 0.22)
##          84) magnet_forearm_z< -294.5 183   32 A (0.83 0.15 0 0.022 0) *
##          85) magnet_forearm_z>=-294.5 2072 1242 C (0.023 0.22 0.4 0.12 0.24)
##          170) magnet_dumbbell_y< 261.5 1052  439 C (0.029 0.18 0.58 0.092 0.12) *
##          171) magnet_dumbbell_y>=261.5 1020  654 E (0.016 0.27 0.21 0.14 0.36)
##          342) magnet_arm_y>=186 449  249 B (0.027 0.45 0.31 0.08 0.14) *
##          343) magnet_arm_y< 186 571  266 E (0.007 0.13 0.13 0.19 0.53) *
##          43) accel_forearm_x< -108.5 916  432 D (0.064 0.079 0.22 0.53 0.11)
##          86) magnet_arm_y>=296.5 269  103 C (0.041 0.12 0.62 0.17 0.056) *
##          87) magnet_arm_y< 296.5 647  209 D (0.074 0.063 0.057 0.68 0.13) *
## 11) magnet_dumbbell_y>=439.5 1507  738 B (0.035 0.51 0.036 0.22 0.2)
## 22) total_accel_dumbbell>=5.5 1076  387 B (0.048 0.64 0.05 0.022 0.24)
## 44) roll_belt>=-0.585 902  213 B (0.058 0.76 0.06 0.027 0.092) *
## 45) roll_belt< -0.585 174    0 E (0 0 0 0 1) *
## 23) total_accel_dumbbell< 5.5 431  128 D (0 0.19 0 0.7 0.11) *
## 3) roll_belt>=130.5 978    8 E (0.0082 0 0 0 0.99) *

```

```

predictions_dt <- predict(model_fit_dt, testing, type = "class")
cm_dt <- confusionMatrix(predictions_dt, testing$classe)
cm_dt

```

```
## Confusion Matrix and Statistics
```

```

##
##          Reference
## Prediction   A    B    C    D    E
##          A 2023  256   36   63   32
##          B   83  958  220  141  188
##          C   55  158  963  107  105
##          D   53   89  100  896  156
##          E   18   57   49   79  961
##

```

```
## Overall Statistics
```

```

##
##          Accuracy : 0.7394
##          95% CI : (0.7295, 0.749)
##          No Information Rate : 0.2845
##          P-Value [Acc > NIR] : < 2.2e-16
##

```

```

##          Kappa : 0.6693
##          McNemar's Test P-Value : < 2.2e-16
##

```

```
## Statistics by Class:
```

```

##
##          Class: A Class: B Class: C Class: D Class: E

```

## Sensitivity	0.9064	0.6311	0.7039	0.6967	0.6664
## Specificity	0.9311	0.9001	0.9344	0.9393	0.9683
## Pos Pred Value	0.8394	0.6025	0.6938	0.6924	0.8256
## Neg Pred Value	0.9616	0.9105	0.9373	0.9405	0.9280
## Prevalence	0.2845	0.1935	0.1744	0.1639	0.1838
## Detection Rate	0.2578	0.1221	0.1227	0.1142	0.1225
## Detection Prevalence	0.3072	0.2027	0.1769	0.1649	0.1484
## Balanced Accuracy	0.9187	0.7656	0.8192	0.8180	0.8174

73.94% overall accuracy. Not the model we want :)

xgboost I am going to try xgboost, another efficient and high performing machine learning algorithm. xgboost works only for numeric data types. So, classe variable in the datasets, which is in factor type, need to be converted to the numeric type.

```
# convert factor type to numeric
classe_factor <- training[, "classe"]
classe_numeric <- classe_factor
class_count <- length(levels(classe_factor))
levels(classe_numeric) = 1:class_count

# Remove the classe column from both training and test data sets
training$classe = NULL
testing$classe = NULL

# Change data type of training and testing to matrix
training_matrix = as.matrix(training)
mode(training_matrix) = "numeric"
testing_matrix = as.matrix(testing)
mode(testing_matrix) = "numeric"
class_outcome <- as.matrix(as.integer(classe_numeric) - 1)
```

Apply the xboost model.

```
model_fit_xgb <-
  xgb.cv(
    param = list("objective" = "multi:softprob", "num_class" = class_count),
    data = training_matrix,
    label = class_outcome,
    nfold = 4,
    nrounds = 100,
    prediction = TRUE,
    verbose = FALSE
  )
predictions_xgb <-
  matrix(
    model_fit_xgb$pred, nrow = length(model_fit_xgb$pred) / class_count, ncol =
      class_count
  )
predictions_xgb <- max.col(predictions_xgb, "last")
confusionMatrix(factor(class_outcome + 1), factor(predictions_xgb))
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    1    2    3    4    5
##           1 3337    7    2    0    2
##           2   18 2255    5    1    0
##           3    0   10 2038    5    1
##           4    0    0   15 1910    5
##           5    0    2    3    7 2153
##
## Overall Statistics
##
##           Accuracy : 0.993
##           95% CI : (0.9913, 0.9944)
##           No Information Rate : 0.2849
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9911
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.9946   0.9916   0.9879   0.9932   0.9963
## Specificity      0.9987   0.9975   0.9984   0.9980   0.9988
## Pos Pred Value   0.9967   0.9895   0.9922   0.9896   0.9945
## Neg Pred Value   0.9979   0.9980   0.9974   0.9987   0.9992
## Prevalence       0.2849   0.1931   0.1752   0.1633   0.1835
## Detection Rate   0.2834   0.1915   0.1731   0.1622   0.1828
## Detection Prevalence 0.2843   0.1935   0.1744   0.1639   0.1838
## Balanced Accuracy 0.9967   0.9946   0.9931   0.9956   0.9975

```

Overall accuracy is 99%+ which means error rate is less than 1%.

xgboost definitely killed it compared to classification tree! Based on CRAN documentation, xgboost package uses parallel computing on a single machine. It also uses efficient linear model solver.